

**University of Central Missouri**  
**Department of Computer Science & Cybersecurity**

**CS5720 Neural network and Deep learning**  
**Spring 2025**

**Home Assignment 2. (Cover Ch 4,5)**

**Student name:**

**VIDHAI DIAMOND GERA**

**700771948**

## Submission Requirements:

- Total Points: 100
- Once finished your assignment push your source code to your repo (GitHub) and explain the work through the ReadMe file properly. Make sure you add your student info in the ReadMe file.
- Submit your GitHub link and video on the BB.
- Comment your code appropriately ***IMPORTANT***.
- Make a simple video about 2 to 3 minutes which includes demonstration of your home assignment and explanation of code snippets.
- Any submission after provided deadline is considered as a late submission.

## Question 1: Cloud Computing for Deep Learning (20 points)

Cloud computing offers significant advantages for deep learning applications.

- (a) Define **elasticity** and **scalability** in the context of cloud computing for deep learning. (10 points)
- (b) Compare **AWS SageMaker**, **Google Vertex AI**, and **Microsoft Azure Machine Learning Studio** in terms of their deep learning capabilities. (10 points)

### Expected Output

Write the definition and comparison for (a) and (b). No code needed.

### Cloud Computing for Deep Learning

#### (a) Elasticity and Scalability in Cloud Computing for Deep Learning

- **Elasticity:** Elasticity refers to the ability of a cloud computing environment to dynamically allocate and deallocate resources based on real-time demand. For deep learning, this means that compute resources such as GPUs and TPUs can be automatically scaled up during training and scaled down when they are no longer needed, optimizing cost-efficiency.
- **Scalability:** Scalability is the capacity of a cloud system to handle increased workload by adding more resources (scale-out) or enhancing existing resources (scale-up). In deep learning, scalability ensures that large models and datasets can be processed efficiently by distributing the workload across multiple computing nodes.

#### (b) Comparison of AWS SageMaker, Google Vertex AI, and Microsoft Azure Machine Learning Studio

<b>Feature</b>	<b>AWS SageMaker</b>	<b>Google Vertex AI</b>	<b>Microsoft Azure Machine Learning Studio</b>
<b>Ease of Use</b>	Provides managed Jupyter notebooks and automated model tuning	Fully managed with AutoML and integrated Jupyter notebooks	Drag-and-drop interface along with Jupyter support
<b>Compute Resources</b>	Supports EC2 instances with GPUs, TPUs, and Auto-scaling	Uses Google Cloud TPUs, GPUs, and AI Platform	Supports Azure VMs with GPU/TPU options
<b>Model Training</b>	Supports distributed training with built-in algorithms	Custom and AutoML model training available	Automated ML and custom model training
<b>Integration</b>	Integrates with AWS ecosystem (S3, Lambda, Redshift)	Connects with Google Cloud services (BigQuery, Dataflow)	Works with Microsoft services (Power BI, Azure Synapse)
<b>Pricing</b>	Pay-as-you-go with spot instance support for cost savings	Per-use pricing with free-tier support	Pay-per-use with reserved capacity options

Each platform offers robust deep learning capabilities, with AWS SageMaker excelling in enterprise-level features, Google Vertex AI providing strong AutoML and TPU integration, and Azure ML Studio offering a balance of low-code and pro-code environments.

## Question 2: Convolution Operations with Different Parameters (20 points)

### Task: Implement Convolution with Different Stride and Padding (10 points)

Write a Python script using **NumPy** and **TensorFlow/Keras** to perform **convolution** on a **5×5 input matrix** using a **3×3 kernel** with varying parameters.

1. Define the following **5×5 input matrix**:

$$\begin{bmatrix} 1 & 2 & 3 & 4 & 5 \\ 6 & 7 & 8 & 9 & 10 \\ 11 & 12 & 13 & 14 & 15 \\ 16 & 17 & 18 & 19 & 20 \\ 21 & 22 & 23 & 24 & 25 \end{bmatrix}$$

2. Define the following **3×3 kernel**:

$$\begin{bmatrix} 0 & 1 & 0 \\ 1 & -4 & 1 \\ 0 & 1 & 0 \end{bmatrix}$$

3. Perform **convolution operations** with:
  - **Stride = 1, Padding = 'VALID'**
  - **Stride = 1, Padding = 'SAME'**
  - **Stride = 2, Padding = 'VALID'**
  - **Stride = 2, Padding = 'SAME'**
4. Print the **output feature maps** for each case.

### Expected Output

Print the output feature maps for

- **Stride = 1, Padding = 'VALID'**
- **Stride = 1, Padding = 'SAME'**
- **Stride = 2, Padding = 'VALID'**

- **Stride = 2, Padding = 'SAME'**

### Question 3: CNN Feature Extraction with Filters and Pooling (30 points)

#### Task 1: Implement Edge Detection Using Convolution (15 points)

Write a Python script using **NumPy** and **OpenCV (cv2)** to apply **edge detection** on an image using a **Sobel filter**.

- Load a grayscale image (you can use any sample image).
- Apply the **Sobel filter** for **edge detection** in the **x-direction** and **y-direction**.
- Display the original image and the filtered images.

Use the following Sobel filters:

$$\text{Sobel X} = \begin{bmatrix} -1 & 0 & 1 \\ -2 & 0 & 2 \\ -1 & 0 & 1 \end{bmatrix}$$

$$\text{Sobel Y} = \begin{bmatrix} -1 & -2 & -1 \\ 0 & 0 & 0 \\ 1 & 2 & 1 \end{bmatrix}$$

#### Task 2: Implement Max Pooling and Average Pooling (15 points)

Write a Python script using **TensorFlow/Keras** to demonstrate **Max Pooling** and **Average Pooling**.

- Create a **random 4x4 matrix** as an input image.
- Apply a **2x2 Max Pooling** operation.
- Apply a **2x2 Average Pooling** operation.
- Print the original matrix, max-pooled matrix, and average-pooled matrix.

### Expected Output

#### Task1: Edge Detection using Sobel Filter

- Display **three images**:
  1. **Original Image**
  2. **Edge detection using Sobel-X**
  3. **Edge detection using Sobel-Y**

#### Task2: Pooling Operations on Random 4x4 Matrix

- Printed **original matrix**, **max pooled matrix**, and **average pooled matrix**.

## Question 4: Implementing and Comparing CNN Architectures (30 points)

### Task 1: Implement AlexNet Architecture (15 points)

Write a **Python script** using **TensorFlow/Keras** to implement a simplified **AlexNet** model with the following layers:

- **Conv2D Layer:** 96 filters, kernel size = (11,11), stride = 4, activation = ReLU
- **MaxPooling Layer:** pool size = (3,3), stride = 2
- **Conv2D Layer:** 256 filters, kernel size = (5,5), activation = ReLU
- **MaxPooling Layer:** pool size = (3,3), stride = 2
- **Conv2D Layer:** 384 filters, kernel size = (3,3), activation = ReLU
- **Conv2D Layer:** 384 filters, kernel size = (3,3), activation = ReLU
- **Conv2D Layer:** 256 filters, kernel size = (3,3), activation = ReLU
- **MaxPooling Layer:** pool size = (3,3), stride = 2
- **Flatten Layer**
- **Fully Connected (Dense) Layer:** 4096 neurons, activation = ReLU
- **Dropout Layer:** 50%
- **Fully Connected (Dense) Layer:** 4096 neurons, activation = ReLU
- **Dropout Layer:** 50%
- **Output Layer:** 10 neurons, activation = Softmax

Print the **model summary** after defining it.

### Task 2: Implement a Residual Block and ResNet (15 points)

Write a Python script to define a **Residual Block** and use it to build a simple **ResNet-like model**.

1. Implement a function `residual_block(input_tensor, filters)` that:
  - Takes an **input tensor**.
  - Applies **two Conv2D layers** (each with 64 filters, kernel size = (3,3), activation = ReLU).
  - Includes a **skip connection** that adds the input tensor to the output before activation.
2. Create a **ResNet model** that:
  - Uses an **initial Conv2D layer** (64 filters, kernel size = (7,7), stride = 2).
  - Applies **two residual blocks**.
  - Ends with a **Flatten layer**, **Dense layer (128 neurons)**, and **Output layer (Softmax)**.

Print the **model summary** after defining it.

## Expected Output

The output should display:

1. The **model summary** for **AlexNet**.
2. The **model summary** for the **ResNet-like model**.